

732A54

Big Data Analytics

Topic: Relational Databases

Huanyu Li

[huanyu.li@liu.se](mailto:huanyu.li@liu.se)

# Literature

- Elmasri, Navathe, Fundamentals of Database Systems, 7<sup>th</sup> edition, Addison Wesley, 2016. Chapters 3-6 and 9; section 7.1.

# Slide

- This slide is based on Patrick's previous slides, and Olaf's 732A57 slides.
- <https://www.ida.liu.se/~732A57/fo/index.en.shtml>

# Outline

- Basic Terminology
- The Database Approach
- The Lecture's Focuses
  - Representing and Storing data
  - EER to Database Schema
  - Querying RDBs using SQL
- Lab RDB

# Basic Terminology

# Most Basic Terminology

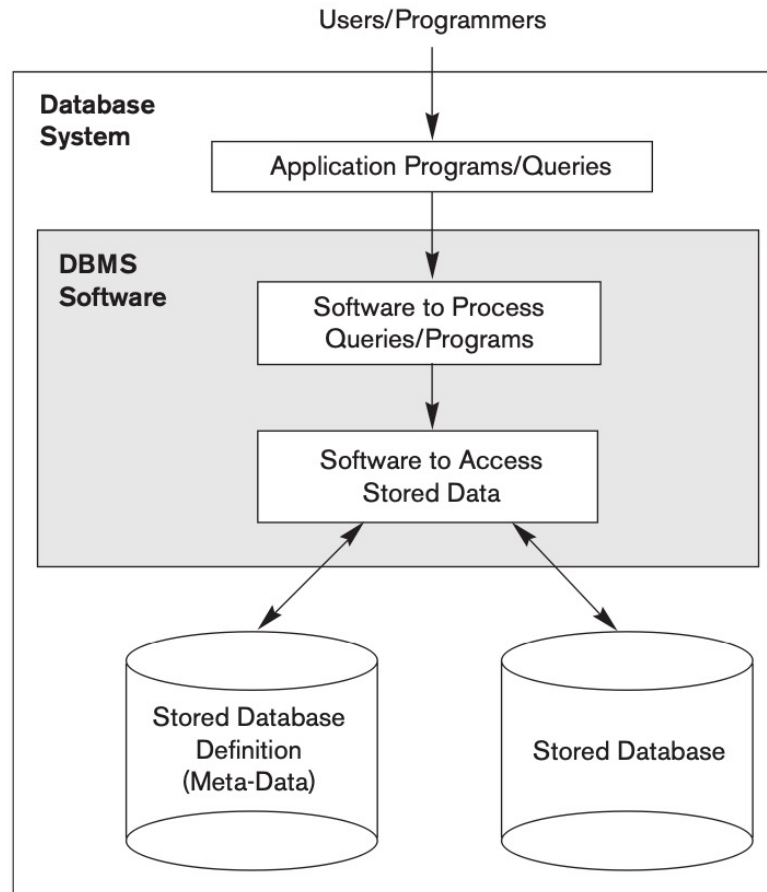
- **Data:** known facts that can be records and that have implicit meaning
- **Database:**
  - collection of related data
  - Represents some aspect of the real world
  - Built for a specific purpose
- Examples of databases (used in every-day life):
  - Amazon.com's product data
  - The Bank's transactions data

# Terminology (cont'd)

- **Database management system (DBMS):**
  - A collection of computer programs
  - Enables users to create and maintain a database (DB)
  - Supports concurrent access to a database by multiple users and programs
  - ...
- Examples of DBMS:
  - Microsoft's SQL Server
  - MySQL
  - Oracle
- **Database System = DB + DBMS**

# Terminology (cont'd)

- **Database System = DB + DBMS**



# Terminology (cont'd)

- **Persons:**

- Database Administrators
- Database Designers
- End Users
  
- System Analysts and Application Programmers
- DBMS designers and implementers
- Tool developers
- Operators and maintenance personnel




# The Database Approach

# Database Approach

- Use a single repository to maintain data that is defined and accessed by various users
- To enable a database approach, we need to design a database system
  - **Database design** focuses on defining the database
  - **Application design** focuses on the programs and interfaces that access the database (out of scope of this lecture)

# Database System Design Process

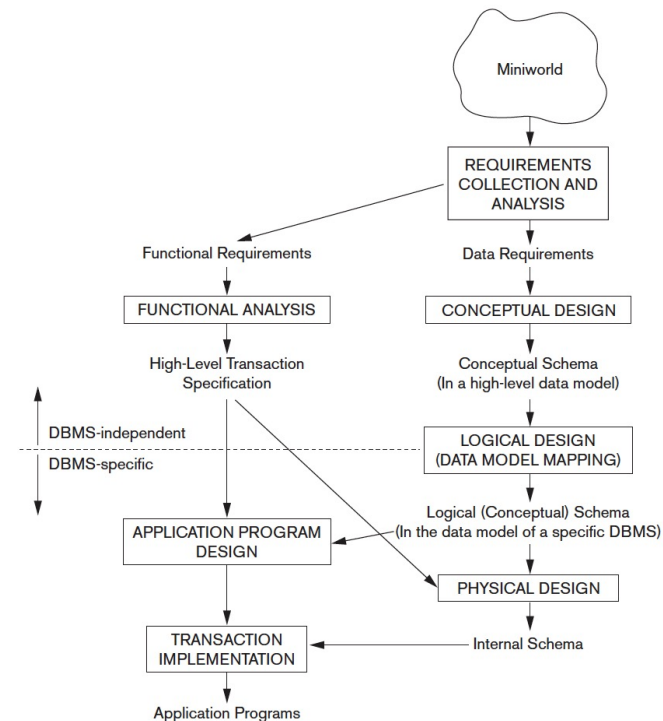
- Database design
  - Requirement specification, analysis
  - Conceptual design
  - Logical design
  - Physical design



# The Lecture's Focuses

# This lecture's focuses

- What information is stored?
- How is the information stored?
  - (high level and low level)
- How is the information accessed?
  - (user level and system level)



# Other issues out of scope of this lecture

- How to optimize performance of a data source?
- How to recover of a data source after crash?
- How to access information from multiple data sources?
- How to allow and control multiple users to access a data source?

# Representing and Storing data

- What information is stored?
- How is the information stored?

# Example

- A taxi company needs to model their activities.
- There are two types of **employees** in the company: **drivers** and **operators**. For drivers it is interesting to know the **date of issue** and **type** of the driving license, and the **date of issue** of the taxi driver's certificate. For all employees it is interesting to know their **personal number**, **address** and the available **phone numbers**.
- The company owns a number of **cars**. For each car there is a need to know its **type**, **year of manufacturing**, **number of places** in the car and **date of the last service**.
- The company wants to have a record of car **trips**. A taxi may be picked on a street or ordered through an **operator** who assigns the order to a certain **driver** and a **car**. **Departure** and **destination addresses** together with **times** should also be recorded.




# ER/EER Model ((Enhanced) Entity Relationship)

- High-level conceptual data model
  - An overview of the database
  - Independent of type of data source
  - Easy to discuss with non-database experts
  - Easy to translate to data model of DBMS
- ER diagram (diagrammatic notion):
  - entities and entity type
  - attributes
  - key attributes
  - relationships and cardinality constraints
  - sub-types (EER)

# Entities and Entity Types

- **Entity**: a ‘thing’ in the real world with an independent existence
- **Attributes**: Properties that describe an entity
- **Entity type**: A collection of entities that have the same set of attributes

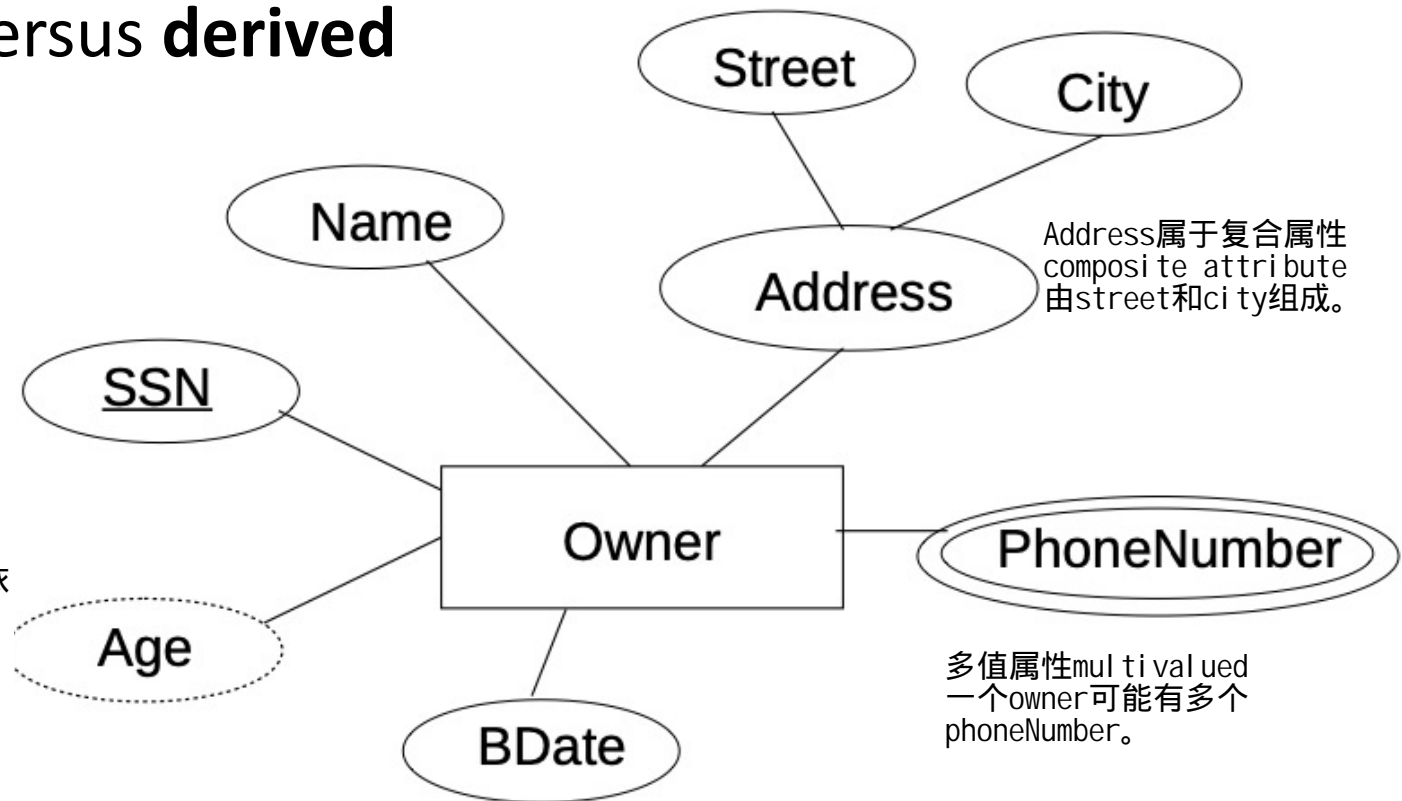


# Attributes

- **simple** versus **composite**
- **single-valued** versus **multivalued**
- **stored** versus **derived**

主键属性  
key attribute  
加下划线  
underline

派生属性derivers attribute  
用虚线椭圆 (Dotted Ellipse)  
, 不直接存储于数据库中, 依  
靠计算推导

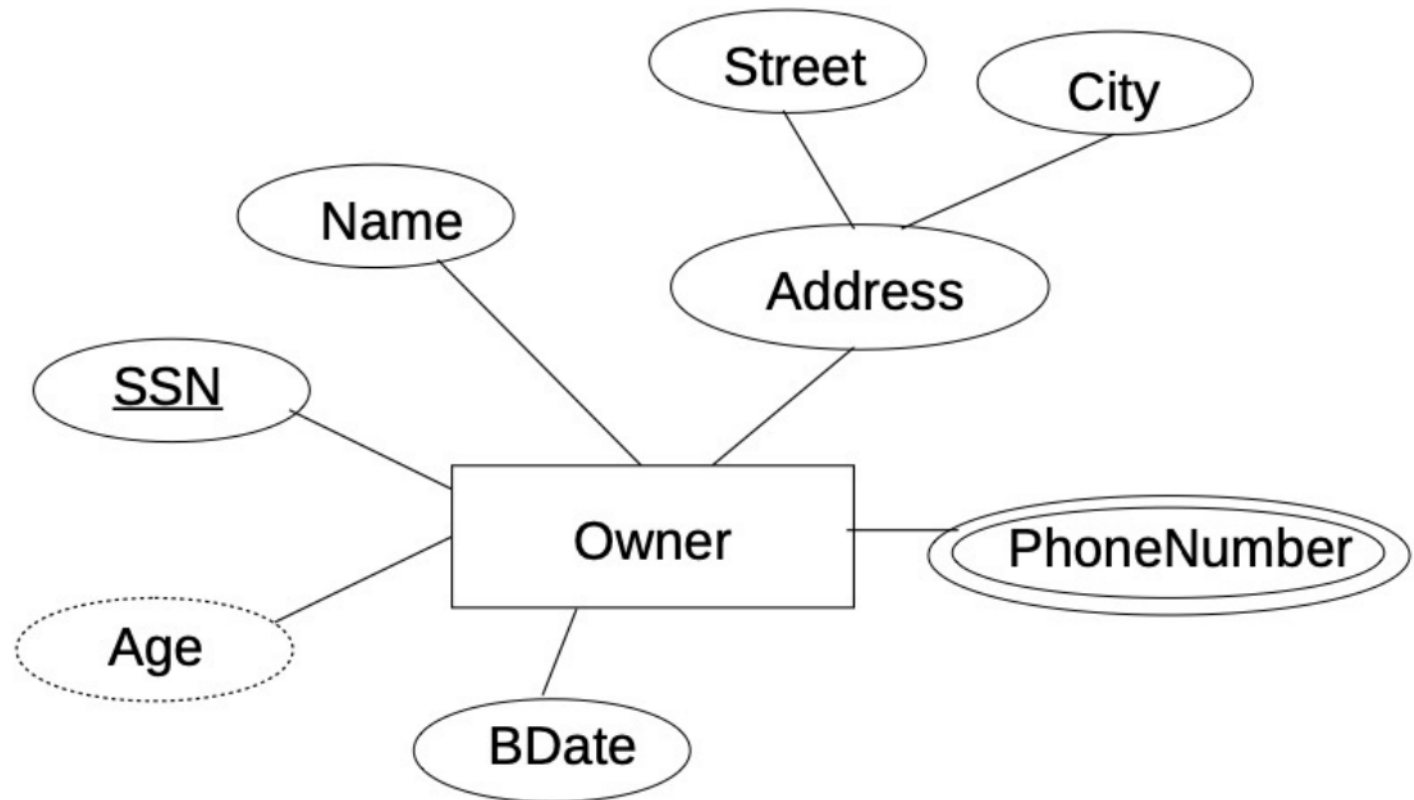


# Constraints on Attributes

- Value sets (domains) of attributes
- Key attributes 主键不能重复

- Domain Constraint: 领域约束

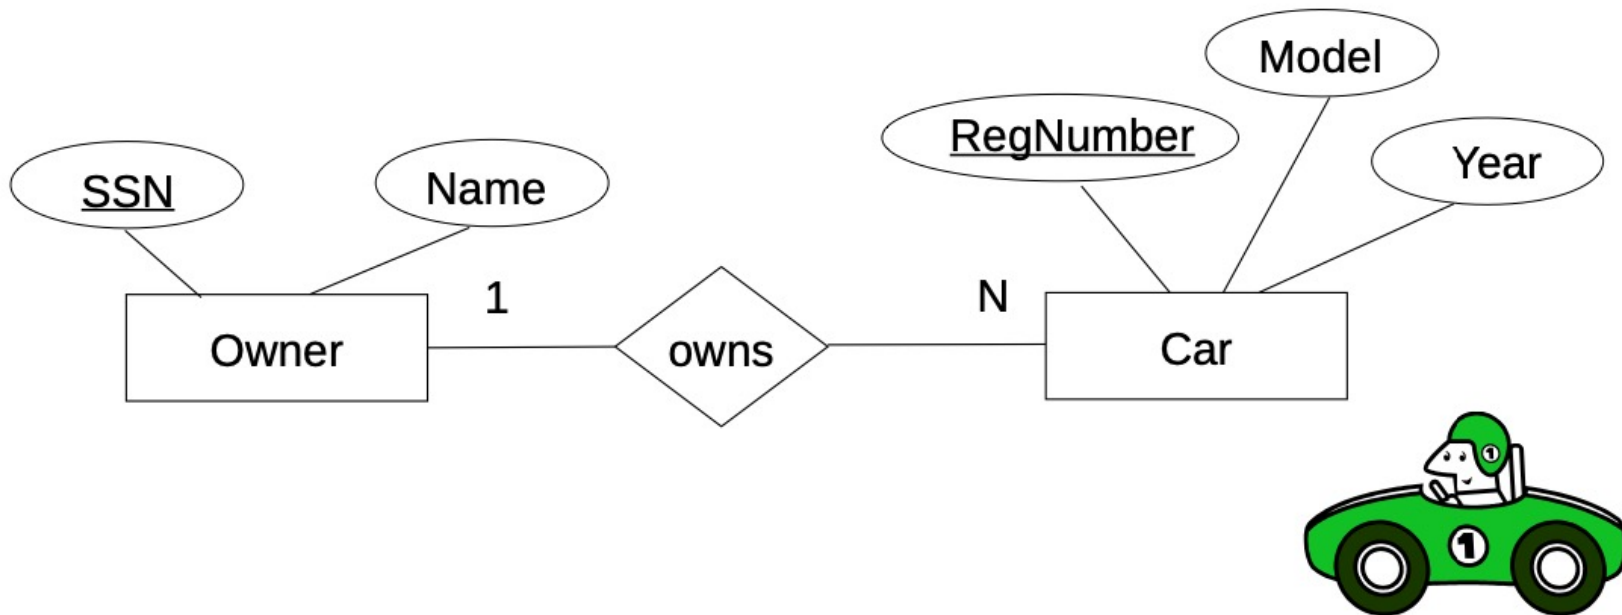
数据类别限制 data type  
数据范围限制 value range  
数据格式限制 format



# Relationship Types

- Relationship type: Association among entity types

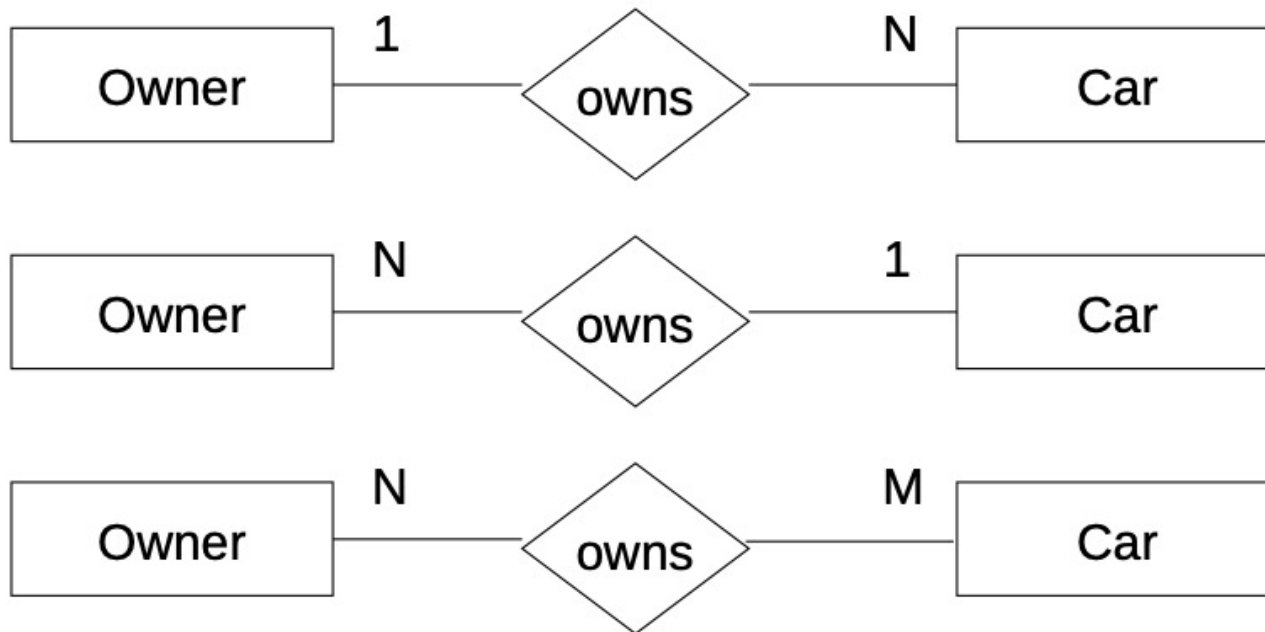
关系类别，用菱形表示关系，数字表示数量关系，下图表示一个人多辆车。关系用来链接两个entity types。



# Constraints on Relationship Types


数量限制，一个实体最多参与多少其他实体的关系

- Cardinality ratio: Maximum number of relationships an entity can participate in
- Possible cardinality ratio:  $1:1$ ,  $1:N$ ,  $N:1$ , and  $N:M$




# Constraints on Relationship Types

- Participation constraint
  - **Total participation:** Every entity participates in at least one relationship with another entity
  - **Partial participation:** Not all entities are involved in the relationship



# Attributes of Relationship Types

”Store information on who owned which car and during which period of time”





# Weak Entity Types

弱实体是完全参与联系的

- Identified by their relationship to a specific entity from another entity type
- Do not have key attributes of their own
  - Only partial key
  - The identifying entity has the rest of the key

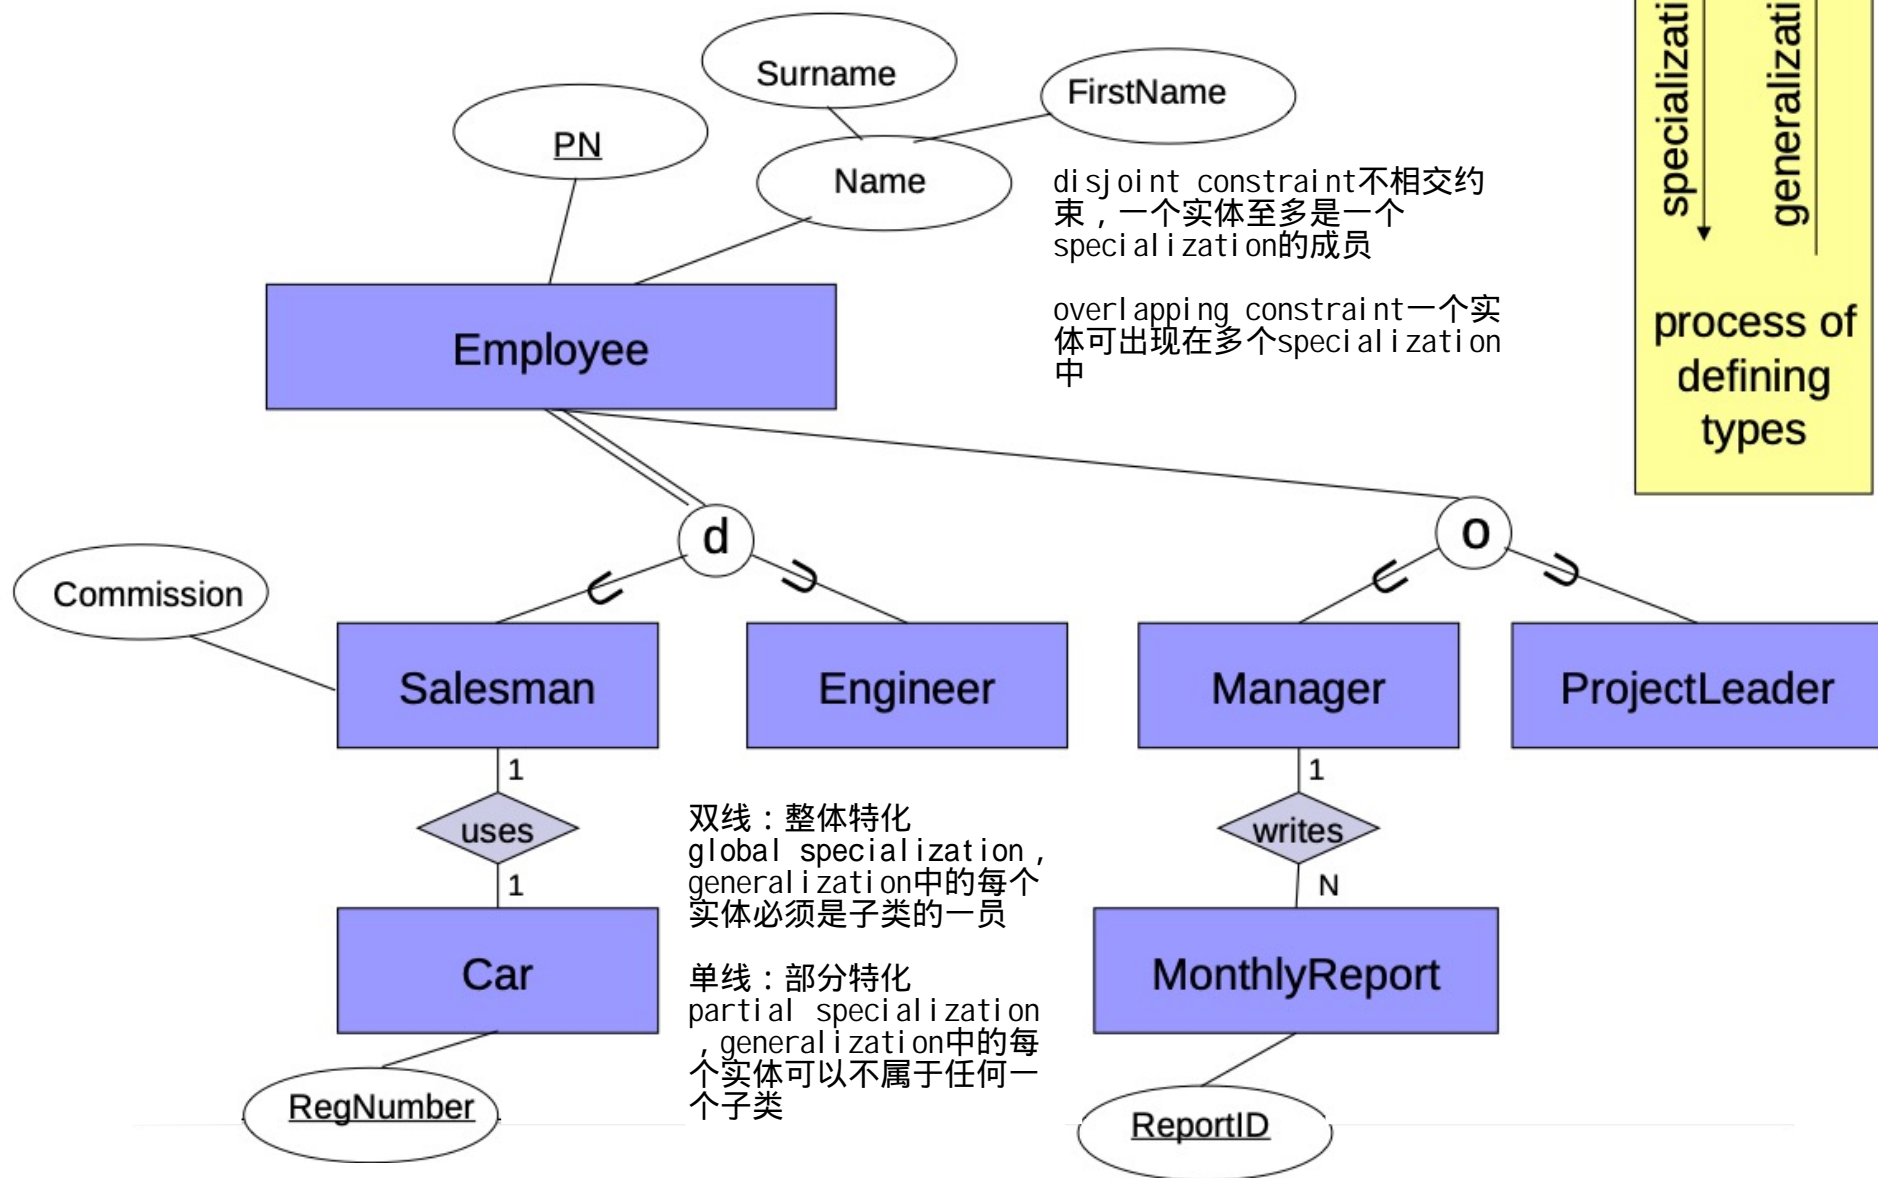
一个强实体和一个弱实体之间的关系用双菱形表示。弱实体在ER图中用双矩形框表示，识别关系用双菱形(Double diamond)表示。部分键属性用虚线表示，部分键(partial key)用于区分不同的弱实体。这里的双线表示全参与，因为room是弱实体，所以room必须依赖于Hotel存在，也就是所有的room都必定对应一个hotel实体，所以room在这个关系中是total的（所以这里要双线表示全参与）。



# Enhanced ER (EER) Model

- Why more? (Enhanced)
  - To support more complex data requirements
  - Example: Only some employees can use a company car, only managers have to write monthly report, but all employees have assigned personal number, salary account and a place in the office
- What more?
  - Specialization / generalization (Subtype / supertype)
  - Union subtypes
  - Attributes and relationship inheritance

# Subtype / Supertype



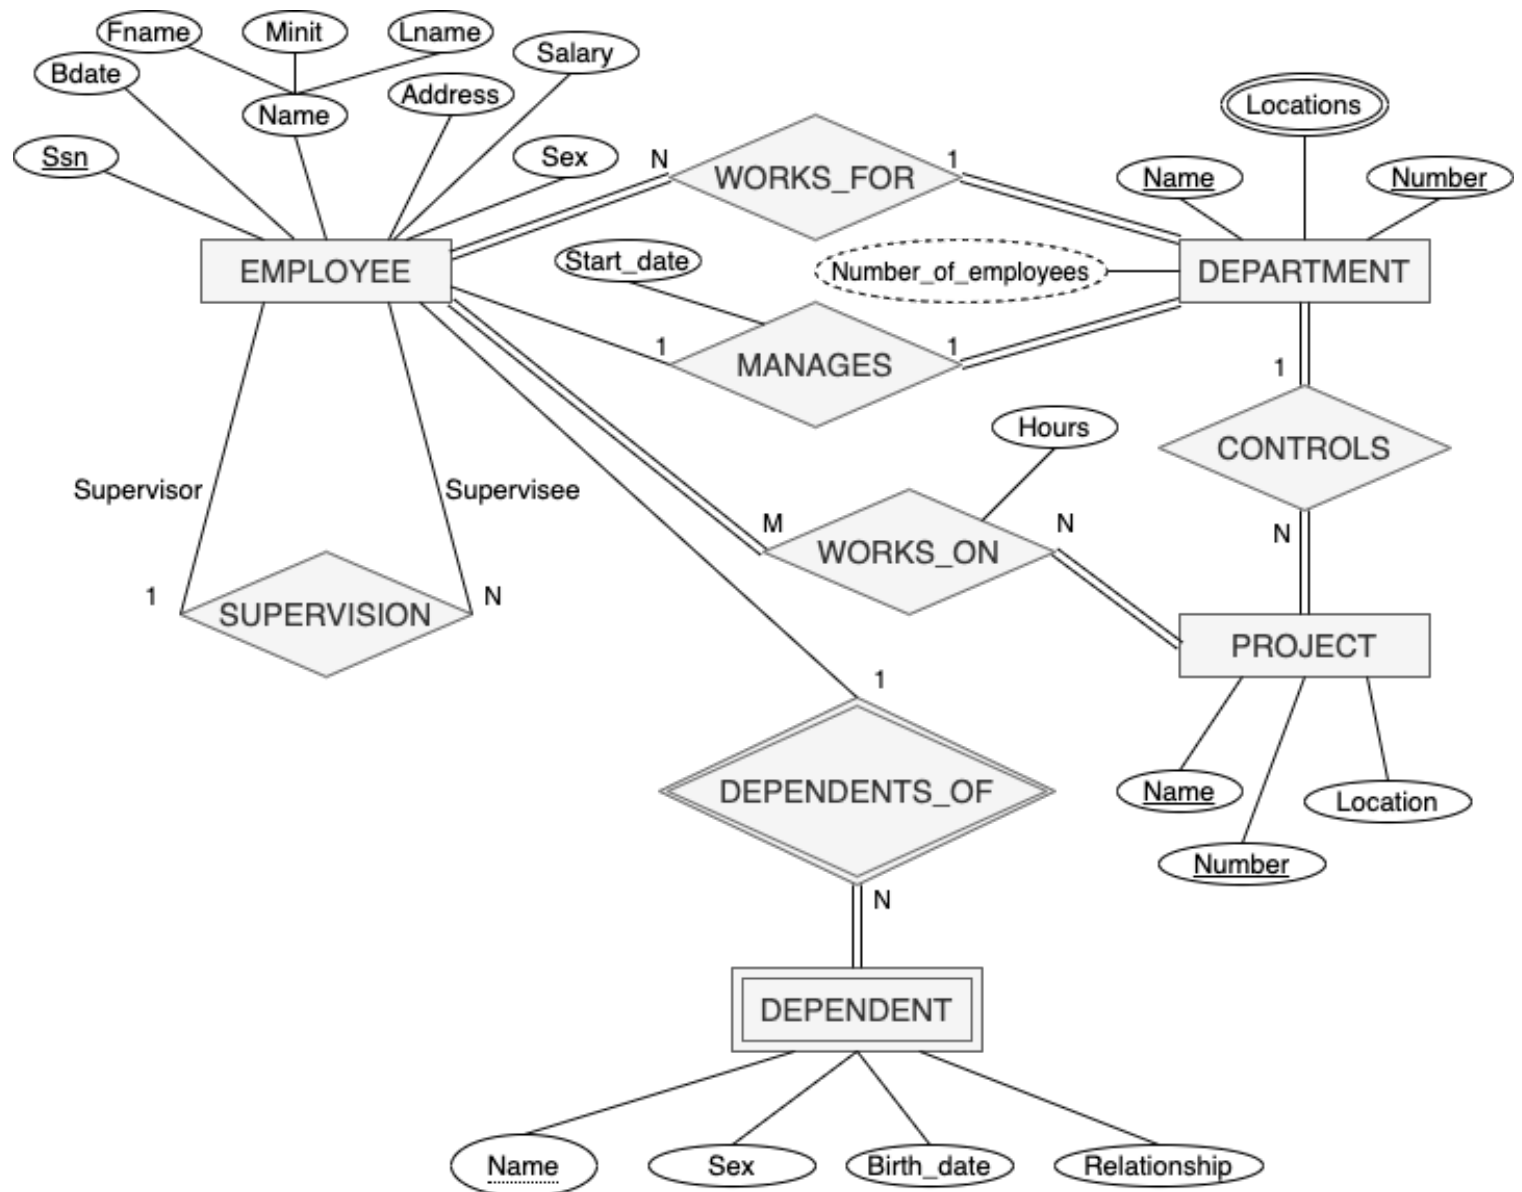
# Exercise: a Miniworld

- The company is organized into departments. Each department has a unique name, a unique number, and a particular employee who manages the department. We keep track of the start date when that employee began managing the department. A department may have several locations.
- A department controls a number of projects, each of which has a unique name, a unique number, and a single location.
- The database will store each employee's name, Social Security number, address, salary, sex (gender), and birth date. An employee is assigned to one department, but may work on several projects, which are not necessarily controlled by the same department. It is required to keep track of the current number of hours per week that an employee works on each project, as well as the direct supervisor of each employee (who is another employee).
- The database will keep track of the dependents of each employee for insurance purposes, including each dependent's first name, sex, birth date, and relationship to the employee.

# Exercise: a Miniworld

- The company is organized into **departments**. Each department has a unique **name**, a unique **number**, and a particular **employee** who **manages** the department. We keep track of the **start date** when that employee began managing the department. A department may have several **locations**.
- A department **controls** a number of **projects**, each of which has a unique **name**, a unique **number**, and a single **location**.
- The database will store each employee's **name**, **Social Security number**, **address**, **salary**, **sex** (gender), and **birth date**. An employee is **assigned** to one department, but may **work on** several projects, which are not necessarily controlled by the same department. It is required to keep track of the current **number of hours per week** that an employee works on each project, as well as the **direct supervisor** of each employee (who is another employee).
- The database will keep track of the **dependents** of each employee for insurance purposes, including each dependent's **first name**, **sex**, **birth date**, and **relationship** to the employee.

# EER



# Storing data with databases

- Relational Databases
  - Model: tables + relation algebra
  - Query language (SQL)
- Object-oriented, NoSQL databases

# Storing data with relational databases

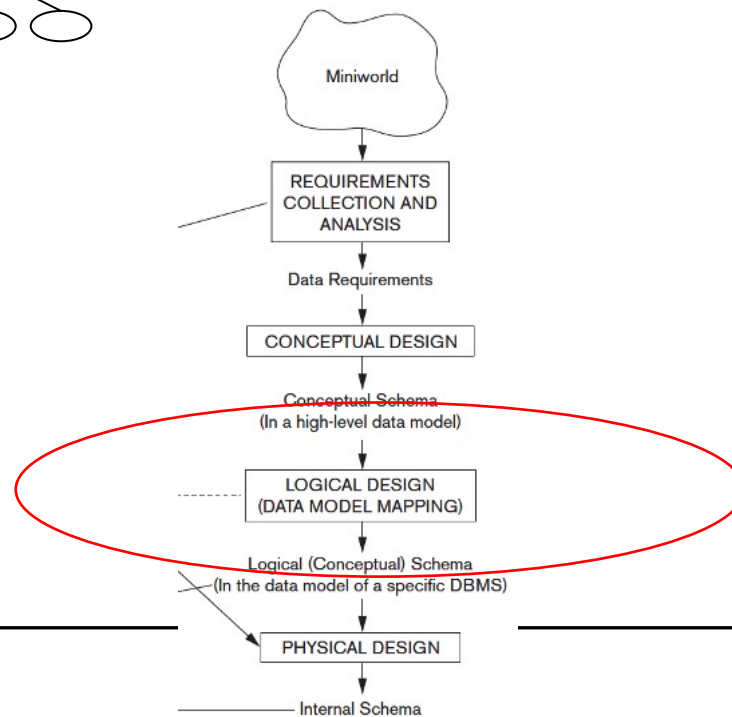



- Relational Database
  - Represent data as a collection of relations
  - Think of a relation as a table of values
  - Each row (tuple) represents a record of related data values
  - Each column (attribute) holds a corresponding value for each row



# EER to Database Schema

图示


# EER to database schema (Logical design)



# EER to database schema

## Step 1: Map Regular Entity Types

- For each regular (strong) entity type  $E$ , create a relation  $R$  that includes all the simple attributes of  $E$



EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary
-------	-------	-------	------------	-------	---------	-----	--------


DEPARTMENT

Dname	<u>Dnumber</u>
-------	----------------

# EER to database schema

## Step 2: Map Weak Entity Types

- For each weak entity type  $W$  with identifying entity type  $E$ , create a relation  $R$  and include all the single-valued attributes of  $W$  and add the primary key attributes from the relation that corresponds to  $E$



DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate
-------------	-----------------------	-----	-------

# EER to database schema

## Step 3: Map Binary 1:1 Relationship Types

- For each binary 1:1 relationship type  $R$ , between  $S$  and  $T$
- Apply one of three possible approaches:
  - 1. Foreign key approach
  - 2. Merged relationship approach
  - 3. Cross-reference or relationship relation approach (Not preferred)

## • 1. Foreign key approach

- Add primary key of one participating relation as foreign key attribute of the other, which will also represent  $R$ 
  - If only on side is total, choose it to represent  $R$
- Declare foreign key attribute as unique
- Add single-valued attributes of relationship type as attribute of  $R$

若S在这个关系中是全参与的，那么就用另一个的主键作为关系R的外部主键。



EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary
-------	-------	-------	------------	-------	---------	-----	--------


DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------



## • 2. Merged relationship approach

- Possible only if both participations are total
- Combine the two relation schemas into one, which will also represent  $R$  直接把两个entity 'cbind' 在一起
- Make one of the primary keys unique instead 随便挑一个的意思
- Add single-valued attributes of relationship type as attributes of  $R$



AB



# EER to database schema

## Step 4: Map 1:N Relationship Types


- For each binary 1:N relationship type  $R$ , between  $S$  and  $T$  (every  $S$  is related to many  $T$ , every  $T$  is related to one  $S$ )
- Foreign key approach
  - Include primary key of entity type  $S$  (1-side) as foreign key in  $T$  (N-side)



# EER to database schema

- Foreign key approach

- Include primary key of entity type *S* (1-side) as foreign key in *T* (N-side)



## EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------




# EER to database schema

## Step 5: Binary M:N and Higher Order Relationship Types

- For each binary M:N relationship type or ternary or higher order relationship type, create a new relation  $S$
- Approach
  - Include primary key of participating entity types as foreign key attributes in  $S$
  - Make all these attributes the primary key of  $S$
  - Include any simple attributes of relationship type in  $S$

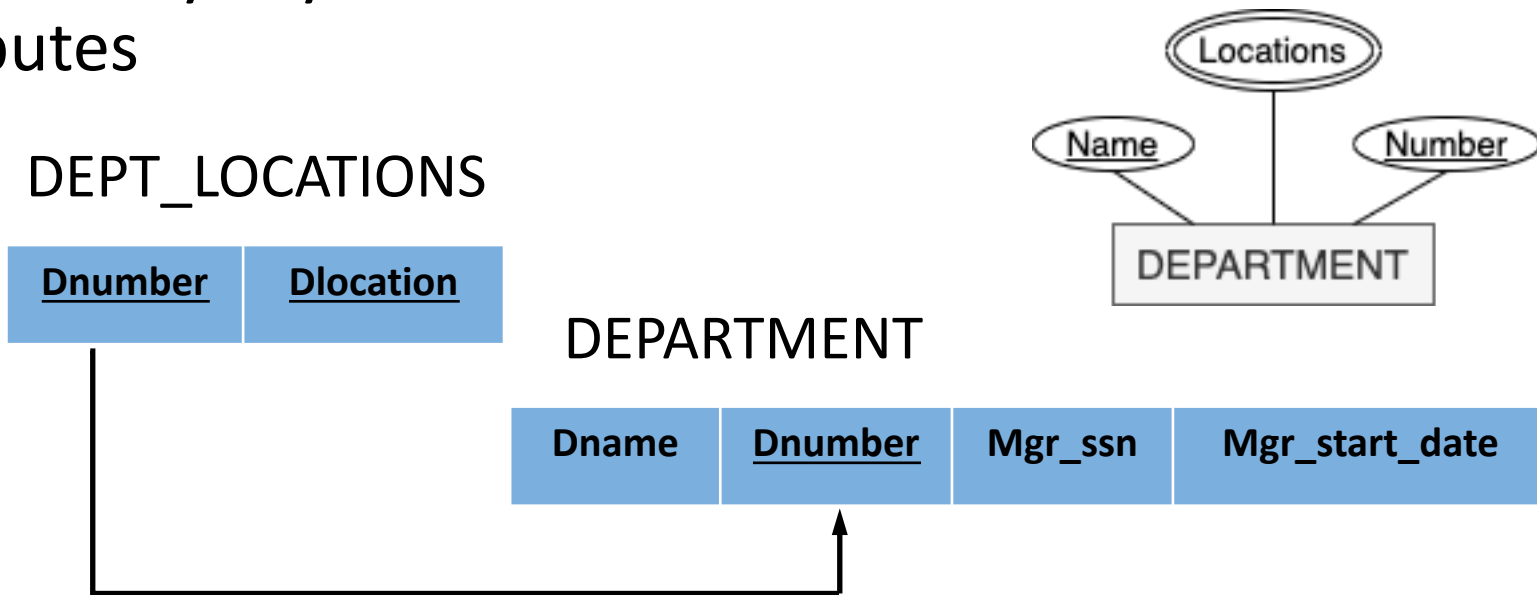
# EER to database schema

- Approach
  - Include primary key of participating entity types as foreign key attributes in *S*
  - Make all these attributes the primary key of *S*
  - Include any simple attributes of relationship type in *S*



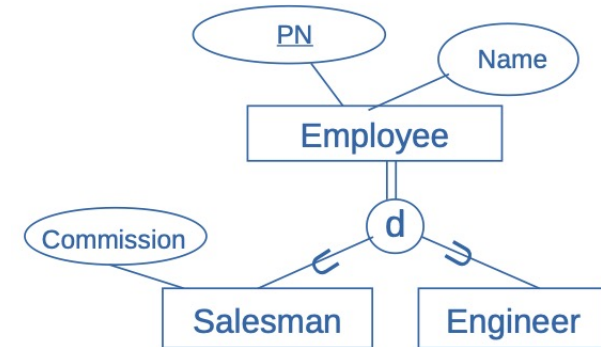

## Step 6: Map Multivalued Attributes

- For each multivalued attribute  $A$ , create a new relation  $R$  that contains an attribute corresponding to  $A$ , and the attribute  $K$ —as a foreign key in  $R$ —of the relation that represents the entity type or relationship type that has  $A$  as a multivalued attribute.
- The primary key of  $R$  is the combination of all its attributes



# Options for mapping specialization/generalization

- For any specialization (total or partial, disjoint or overlapping)
  - 1. Separate relation per subtype*





- 2. Single relation with Boolean type attributes for every subtype, add all attributes of all subtypes*



# Options for mapping specialization/generalization

- For total specializations only
  - *Separate relation per subtype only, overlapping subtypes will result in multiple tuples per entity*
  - *Overlapping subtypes will result in multiple tuples per entity*

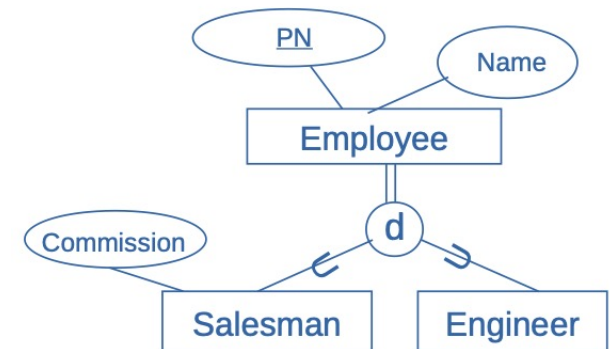


# EER to database schema

## Options for mapping specialization/generalization

- For disjoint specializations only 不相交特化
  - *Single relation with one type attribute*
  - *Type or discriminating attribute indicates subtypes of tuple*
  - *Might require many NULL values if several specific attributes exist in subtypes*

不相交特化有些tuple属于超类，所以子类的一些属性会没有，而且属于子类A的可能没有子类B的一些属性



# Querying RDBs using SQL

-- How is the information accessed?



# SQL

- Developed by IBM Research
- Used in many database systems
- Statements for data definitions, queries and updates
- Tutorial:
  - <https://www.w3schools.com/sql/>

# Relational Schema

## EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPER_SSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

DNAME	<u>DNUMBER</u>	MGR_SSN	MGR_START_DATE
-------	----------------	---------	----------------

## DEPT\_LOCATIONS

<u>DNUMBER</u>	<u>DLOCATION</u>
----------------	------------------

# Relational Schema (cont'd)

## PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

## WORKS\_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

## DEPENDENT

<u>ESSN</u>	<u>DEPARTMENT_NAME</u>	SEX	BDATE
-------------	------------------------	-----	-------

# SQL – Creating Tables

```
CREATE TABLE <tablename> (  
    <colname> <datatype> [<constraint>],  
    ...,  
    [<constraint>],  
    ...  
);
```

- Data types: integer, decimal, number, varchar, char, etc.
- Constraints: not null, primary key, foreign key, unique, etc.

# SQL – Creating Tables (Example)

```
CREATE TABLE WORKS_ON (  
    ESSN      integer,  
    PNO       integer,  
    HOURS     decimal(3, 1),  
  
    constraint pk_workson primary key (ESSN, PNO),  
    constraint fk_works_emp FOREIGN KEY (ESSN) references EMPLOYEE(SSN),  
    constraint fk_works_proj FOREIGN KEY (PNO) references PROJECT(PNUMBER)  
);
```

# SQL – Modifying Table Definitions(Example)

- Add, delete, and modify columns and constraints

```
ALTER TABLE EMPLOYEE ADD COLUMN JOB VARCHAR(12);
```

```
ALTER TABLE EMPLOYEE DROP COLUMN ADDRESS;
```

```
ALTER TABLE WORKS_ON DROP FOREIGN KEY fk_works_emp;
```

```
ALTER TABLE WORKS_ON ADD CONSTRAINT fk_works_emp  
    FOREIGN KEY (ESSN) REFERENCES EMPLOYEE(SSN);
```

- Delete a table and its definition

```
DROP TABLE EMPLOYEE;
```

# SQL Retrieval Queries - Syntax

- Using a SELECT statement

**SELECT** *<attribute-list>*

**FROM** *<table-list>*

**[ WHERE** *<condition>* **] ;**

- *<attribute-list>* is a list of column names (or expressions) whose values are to be retrieved
- *<table-list>* is a list of table names required to process the query
- *<condition>* is a Boolean expression that identifies the tuples to be retrieved by the query (if no WHERE clause, all tuples to be retrieved)

# SQL Retrieval Queries -Examples

- Q1: List SSN for all employees.

**SELECT SSN**

**FROM EMPLOYEE;**

SSN

---

123456789

333445555

999887777

987654321

666884444

453453453

987987987

888665555



# SQL Retrieval Queries -Examples

- Q2: List birth date and address for all employees whose name is 'John B. Smith'.

**SELECT** BDATE, ADDRESS

**FROM** EMPLOYEE

**WHERE** FNAME = 'John'

**AND** MINIT = 'B'

**AND** LNAME = 'Smith';

BDATE

ADDRESS

---

1965-01-09 731	Fondren, Houston, TX
----------------	----------------------

---

# SQL Retrieval Queries -Examples

- Q3: List all information about the employees of department 5.

```
SELECT FNAME, MINIT, LNAME,SSN, BDATE,  
        ADDRESS, SEX, SALARY, SUPERSSN, DNO  
FROM EMPLOYEE  
WHERE DNO = 5;
```

# SQL Retrieval Queries -Examples

- Q4: List all employees that work at the research department.

```
SELECT *  
FROM EMPLOYEE, DEPARTMENT  
WHERE DNAME= 'Research'  
AND DNUMBER = DNO;
```

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPER_SSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGR_SSN	MGR_START_DATE
-------	----------------	---------	----------------

# SQL Retrieval Queries -Examples

- Q5: List project number, department number and the name and address of the director of the department for all projects that are located in Stafford.

```
SELECT PNUMBER, DNUM, LNAME, ADDRESS  
FROM PROJECT, DEPARTMENT, EMPLOYEE  
WHERE PLOCATION = `Stafford'  
        AND DNUMBER = DNUM  
        AND SSN = MGRSSN;
```

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

# SQL Retrieval Queries -Examples

```
SELECT PROJECT.PNUMBER, PROJECT.DNUM,  
        EMPLOYEE.LNAME, EMPLOYEE.ADDRESS  
FROM PROJECT, DEPARTMENT, EMPLOYEE  
WHERE PROJECT.PLOCATION = `Stafford'  
        AND DEPARTMENT.DNUMBER = PROJECT.DNUM  
        AND EMPLOYEE.SSN = DEPARTMENT.MGRSSN;
```

# SQL Retrieval Queries -Examples

- Q6: List first and last name for all employees together with first and last names of their bosses.

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM EMPLOYEE E, EMPLOYEE S  
WHERE E.SUPERSSN = S.SSN;
```

# SQL Retrieval Queries -Examples

- Q7: List all salaries.

```
SELECT SALARY  
FROM EMPLOYEE;
```

# Tables as Sets

- SQL considers a table as a multi-set (bag), i.e. tuples can occur more than once in a table.
- Why?
  - Removing duplicates is expensive.
  - User may want information about duplicates.
  - Aggregation operators.

<b><i>SALARY</i></b>
30000
40000
25000
43000
38000
25000
25000
55000



# SQL Retrieval Queries -Examples

- Q8: List all salaries without duplicates.

```
SELECT DISTINCT SALARY  
FROM EMPLOYEE;
```

***SALARY***

---

30000

40000

25000

43000

38000

55000

# SQL Retrieval Queries -Examples

- Q9: List all project numbers for projects in which an employee with name Smith works **or** where the leader of the department to which the project belongs is called Smith.

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPER_SSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT  
T

DNAME	<u>DNUMBER</u>	MGR_SSN	MGR_START_DATE
-------	----------------	---------	----------------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS\_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

# SQL Retrieval Queries -Examples

```
(SELECT DISTINCT PNUMBER  
FROM PROJECT, DEPARTMENT, EMPLOYEE  
WHERE DNUM = DNUMBER AND  
MGRSSN = SSN AND LNAME = 'Smith')  
UNION  
(SELECT DISTINCT PNUMBER  
FROM PROJECT, WORKS_ON, EMPLOYEE  
WHERE PNO = PNUMBER AND  
ESSN = SSN AND LNAME = 'Smith');
```

# SQL Retrieval Queries -Examples

- Q10: List all employees that live in Houston.

```
SELECT FNAME, LNAME  
FROM EMPLOYEE  
WHERE ADDRESS LIKE '%Houston%';
```

# SQL Retrieval Queries -Examples

- Q11: List names for all employees that are born in the 1950's.

```
SELECT FNAME, LNAME  
FROM EMPLOYEE  
WHERE BDATE LIKE '__ _ 5%';
```

# SQL Retrieval Queries -Examples

- Q12: List names and salaries for all employees that work with ProductX in case they would receive a raise of 10%.

```
SELECT FNAME, LNAME, 1.1 * SALARY  
FROM EMPLOYEE, WORKS_ON, PROJECT  
WHERE SSN = ESSN  
      AND PNO = PNUMBER  
      AND PNAME = `PRODUCTX`;
```

# SQL Retrieval Queries -Examples

- Q13: List all employees in department 5 with a salary between 30,000\$ and 40,000\$.

```
SELECT *  
FROM EMPLOYEE  
WHERE DNO = 5 AND  
(SALARY BETWEEN 30000 AND 40000);
```

# SQL Retrieval Queries -Examples

- Q14: List all employees and the projects they work with sorted with respect to department and within the department sorted alphabetically with respect to last name and then first name.

```
SELECT DNAME, LNAME, FNAME, PNAME  
FROM DEPARTMENT, EMPLOYEE, PROJECT, WORKS_ON  
WHERE PNO = PNUMBER AND SSN = ESSN AND DNO =  
DNUMBER  
ORDER BY DNAME, LNAME, FNAME;
```



# SQL Retrieval Queries -Examples

```
SELECT DNAME, LNAME, FNAME, PNAME  
FROM DEPARTMENT, EMPLOYEE, PROJECT, WORKS_ON  
WHERE PNO = PNUMBER AND SSN = ESSN AND DNO =  
    DNUMBER  
ORDER BY DNAME DESC, LNAME ASC, FNAME ASC;
```

# SQL Retrieval Queries -Examples

- Q15: List SSN for all employees that work with the same project at the same times as the person with SSN '123456789' (John Smith).

```
SELECT ESSN  
FROM WORKS_ON  
WHERE (PNO, HOURS) IN  
      (SELECT PNO, HOURS  
      FROM WORKS_ON  
      WHERE ESSN = '123456789');
```

# SQL Retrieval Queries -Examples

```
SELECT E.ESSN  
FROM WORKS_ON E, WORKS_ON JS  
WHERE JS.ESSN = '123456789'  
      AND E.PNO = JS.PNO  
      AND E.HOURS = JS.HOURS;
```

# SQL Retrieval Queries -Examples

- Q16: List all employees whose salary is higher than the salaries of the employees who work at department 5.

```
SELECT LNAME, FNAME  
FROM EMPLOYEE  
WHERE SALARY > ALL  
    (SELECT SALARY  
    FROM EMPLOYEE  
    WHERE DNO = 5);
```

# SQL Retrieval Queries -Examples

- Q17: List all employees whose salary is higher than the salary of some employee who works at department 5.

```
SELECT LNAME, FNAME  
FROM EMPLOYEE  
WHERE SALARY > SOME  
    (SELECT SALARY  
    FROM EMPLOYEE  
    WHERE DNO = 5);
```

# SQL Retrieval Queries -Examples

- Q18: List all employees that do not have a relative at the company.

```
SELECT LNAME, FNAME  
FROM EMPLOYEE  
WHERE NOT EXISTS  
  (SELECT *  
   FROM DEPENDENT  
   WHERE SSN = ESSN);
```

# SQL Retrieval Queries -Examples

- Q19: List all department managers that have at least one relative at the company.

```
SELECT LNAME, FNAME  
FROM EMPLOYEE  
WHERE EXISTS  
    (SELECT *  
     FROM DEPARTMENT  
     WHERE SSN = MGRSSN)  
 AND EXISTS  
    (SELECT *  
     FROM DEPENDENT  
     WHERE SSN = ESSN);
```

# SQL Retrieval Queries -Examples

- Q20: List SSN for all employees that work with project 1, 2 or 3.

```
SELECT DISTINCT ESSN  
FROM WORKS_ON  
WHERE PNO IN (1, 2, 3);
```



# SQL Retrieval Queries -Examples

- Q21: List all employees that do not have a boss.

```
SELECT FNAME, LNAME  
FROM EMPLOYEE  
WHERE SUPERSSN IS NULL;
```

# SQL Retrieval Queries -Examples

- Q22: List the sum, the highest, lowest and average of the salaries of the employees of the research department.

```
SELECT SUM(SALARY), MAX(SALARY), MIN(SALARY),  
        AVG(SALARY)  
FROM EMPLOYEE, DEPARTMENT  
WHERE DNAME = 'Research'  
      AND DNO = DNUMBER;
```

# SQL Retrieval Queries -Examples

- Q23: List the number of employees.

```
SELECT COUNT(*)  
FROM EMPLOYEE;
```

## SQL Retrieval Queries -Examples

- Q24: List for each department the department number, the number of employees and the average salary.

```
SELECT DNO, COUNT(*), AVG(SALARY)  
FROM EMPLOYEE  
GROUP BY DNO;
```

DNO	COUNT(*)	AVG_SALARY
5	4	33250
4	3	31000
1	1	55000

# SQL Retrieval Queries -Examples

- Q25: List for each project the project number, project name and the number of employees that work with the project.

```
SELECT PNUMBER, PNAME, COUNT(*)  
FROM PROJECT, WORKS_ON  
WHERE PNUMBER = PNO  
GROUP BY PNUMBER, PNAME;
```

# SQL Retrieval Queries -Examples

- Q26: List for each project with at least 2 employees the project number, project name and number of employees that work with the project.

```
SELECT PNUMBER, PNAME, COUNT(*)  
FROM PROJECT, WORKS_ON  
WHERE PNUMBER = PNO  
GROUP BY PNUMBER, PNAME  
HAVING COUNT(*) > 1;
```

# Extended SELECT Syntax

- **SELECT** *<attribute-list and function-list>*  
**FROM** *<table-list>*  
[ **WHERE** *<condition>* ]  
[ **GROUP BY** *<grouping attribute-list>*]  
[ **HAVING** *<group condition>* ]  
[ **ORDER BY** *<attribute-list>* ];


# Lab RDB



# Lab RDB

- <https://www.ida.liu.se/~732A54/lab/rdb/index.en.shtml>
- The Jonson Brothers' database
  - The lab exercises are based on a database that is used for the business of the Jonson Brothers. The Jonson Brothers is a retail company with department stores in many major US cities. The company has a large number of employees and sells a varied line of products. The company consists of a number of stores that contain a number of departments. The company has employees, who (among other things) sell items at the different stores. Sales are registered in the sale and debit tables. Items are bought from various suppliers, who also supply parts for the company's computer equipment. Deliveries of computer parts are registered in the supply table.
- Answer questions 1 to 14

# The Jonson Brothers' ER-Diagram



# Tables

## *jbemployee*

An employee is identified by an id and described by name, salary, birthyear and startyear. The id of the manager of each employee is also supplied. A null value means that the employee has no manager.

id	name	salary	manager	birthyear	startyear
157	Jones, Tim	12000	199	1940	1960
1110	Smith, Paul	6000	33	1952	1973
35	Evans, Michael	5000	32	1952	1974
129	Thomas, Tom	10000	199	1941	1962
13	Edwards, Peter	9000	199	1928	1958
215	Collins, Joanne	7000	10	1950	1971
55	James, Mary	12000	199	1920	1969
26	Thompson, Bob	13000	199	1930	1970
98	Williams, Judy	9000	199	1935	1969
32	Smythe, Carol	9050	199	1929	1967
33	Hayes, Evelyn	10100	199	1931	1963
199	Bullock, J.D.	27000	NULL	1920	1920
4901	Bailey, Chas M.	8377	32	1956	1975
843	Schmidt, Herman	11204	26	1936	1956
2398	Wallace, Maggie J.	7880	26	1940	1959
1639	Choy, Wanda	11160	55	1947	1970
5119	Bono, Sonny	13621	55	1939	1963
37	Raveen, Lemont	11985	26	1950	1974
5219	Schwarz, Jason B.	13374	33	1944	1959
1523	Zugnoni, Arthur A.	19868	129	1928	1949
430	Brunet, Paul C.	17674	129	1938	1959
994	Iwano, Masahiro	15641	129	1944	1970
1330	Onstad, Richard	8779	13	1952	1971
10	Ross, Stanley	15908	199	1927	1945
11	Ross, Stuart	12067	NULL	1931	1932

## *jbdept*

A department is identified by an id and described by its name as well as its which store and floor it belongs to. The employee id of the manager of the department is also supplied.

id	name	store	floor	manager
35	Book	5	1	55
10	Candy	5	1	13
19	Furniture	7	4	26
20	MajorAppliances	7	4	26
14	Jewelry	8	1	33
43	Children's	8	2	32
65	Junior's	7	3	37
58	Men's	7	2	129
60	Sportswear	5	1	10
99	Giftwrap	5	1	98
1	Bargain	5	0	37
26	Linens	7	3	157
63	Women's	7	3	32
49	Toys	8	2	35
70	Women's	5	1	10
73	Children's	5	1	10
34	Stationary	5	1	33
47	JuniorMiss	7	2	129
28	Women's	8	2	32

# Getting started

- <https://www.ida.liu.se/~732A54/lab/rdb/index.en.shtml>
- MySQL

## RDB Lab sessions

03-31 08:15-10:00, SU00-01, Group A, Huanyu Li

03-31 08:15-10:00, SU02-03, Group B, Sebastian Ferrada

04-05 15:15-17:00, SU00-01, Group A, Huanyu Li

04-05 15:15-17:00, SU02-03, Group B, Sebastian Ferrada

## Submission

<https://gitlab.liu.se/olaha93/bigdata>