

Computer Lab 2

Computational Statistics

Linköpings Universitet, IDA, Statistik

November 9, 2021

Course code and name:	732A90 Computational Statistics
Lab session:	12.11, 8-10
Submission deadline:	16.11, 23:59
Resubmission deadlines:	resubmission 1: 30.11, 23:59; resubmission 2 for labs 1-4: 14.1
Seminar:	Seminar 1 (first part) on 1.12
Teachers:	Maryna Prus, Filip Ekström, Joel Oskarsson, Martynas Lukosevicius, Shashi Nagarajan, Yifan Ding
Instructions:	This computer laboratory is a part of the examination Create a group report (in English) on the solutions to the lab as a .PDF file. All R codes should be included as an appendix into your report. In the report reference all consulted sources and disclose all collaborations. The report should be handed in via LISAM (or alternatively in case of problems e-mailed to your teacher - see file "lab groups" on lisam).

Exercises originally developed by Krzysztof Bartoszek

Question 1: Optimizing parameters

Finding the minimum or maximum of a function is usually presented as a goal in itself. Here you are asked to use the function `optim()` to create a procedure to approximate another function, through so-called parabolic interpolation. For this exercise let $f(x)$ be a continuous function on the interval $[0, 1]$ and let $x_0, x_1, x_2 \in [0, 1]$ be such that $f(x_1) < f(x_0), f(x_2)$. We will approximate the function $f(x)$ with a function that is piecewise $\tilde{f}(x) = a_0 + a_1x + a_2x^2$, i.e. a piecewise quadratic function.

1. Write a function that uses `optim()` and finds values of (a_0, a_1, a_2) for which \tilde{f} interpolates f at user provided points x_0, x_1, x_2 . Interpolate means $f(x_0) = \tilde{f}(x_0)$, $f(x_1) = \tilde{f}(x_1)$ and $f(x_2) = \tilde{f}(x_2)$. `optim()` should minimize the squared error, i.e. find (a_0, a_1, a_2) that make $(f(x_0) - \tilde{f}(x_0))^2 + (f(x_1) - \tilde{f}(x_1))^2 + (f(x_2) - \tilde{f}(x_2))^2$ as small as possible.
2. Now construct a function that approximates a function defined on the interval $[0, 1]$. Your function should take as a parameter the number of equal-sized intervals that $[0, 1]$ is to be divided into and the function to approximate. The target function is known at the ends of the interval and also at the mid-point of the interval. Independently on each subinterval you should approximate the target function using the **parabolic interpolater** implemented in the previous part i.e. use the parabolic interpolater to find a_0, a_1, a_2 for each subinterval.
3. Apply your function from the previous item to $f_1(x) = -x(1-x)$ and $f_2(x) = -x \sin(10\pi x)$. Plot $f_1(\cdot)$, $\tilde{f}_1(\cdot)$ and $f_2(\cdot)$, $\tilde{f}_2(\cdot)$. How did your piecewise-parabolic interpolater fare? Explain what you observe. Take the number of subintervals to be at least 100.

Question 2: Maximizing likelihood

The file `data.RData` contains a sample from normal distribution with some parameters μ, σ . For this question read `?optim` in detail.

1. Load the data to R environment.
2. Write down the log-likelihood function for 100 observations and derive maximum likelihood estimators for μ, σ analytically by setting partial derivatives to zero. Use the derived formulae to obtain parameter estimates for the loaded data.
3. Optimize the minus log-likelihood function with initial parameters $\mu = 0, \sigma = 1$. Try both Conjugate Gradient method (described in the presentation handout) and BFGS (discussed in the lecture) algorithm with gradient specified and without. Why it is a bad idea to maximize likelihood rather than maximizing log-likelihood?
4. Did the algorithms converge in all cases? What were the optimal values of parameters and how many function and gradient evaluations were required for algorithms to converge? Which settings would you recommend?