

Model Selection and Hypothesis Testing

732A90

Computational Statistics

Maryna Prus


(maryna.prus@liu.se)

Slides originally by Krzysztof Bartoszek

November 23, 2021

Department of Computer and Information Science
Linköping University

Model Selection



Tools for model selection

- Comparing different models
 - Hypothesis testing
- Uncertainty estimation
 - Confidence intervals

Hypothesis Testing

❶ State

H_0 - null hypothesis (e.g. no difference between groups)

H_1 - alternative (e.g. there is difference)

❷ Observe data

❸ Calculate test statistic

❹ Under H_0 test statistic has “known” distribution

❺ Decision:

Is value of test statistic *surprising* (in *critical region*)?

If so reject H_0 in favour of H_1

❻ Two-sided vs. one-sided tests

Hypothesis Testing: Example

x_1, \dots, x_n - data from $N(\mu, \sigma^2)$

Hypotheses:

$$H_0 : \mu = 4$$

$$H_1 : \mu \neq 4$$

Test statistic

$$T(x) = \frac{\bar{x} - \mu}{s/\sqrt{n}} \sim t_{n-1}$$

```
tx<-(mean(x)-4)/(sqrt(var(x)/length(x)))  
t0<-qt(0.975,df=length(x)-1)  
(tx>t0) || (tx<(-t0)) ## reject if TRUE
```

Hypothesis Testing: Power


Type II error - probability not to reject H_0 if it is "false"

Power = 1 - Type II error

- ability to correctly identify
surprise,
i.e. indicate H_1

Computing power

- Analytically ?
- Generate data samples that satisfy H_1
Compute percent of correct rejections



Source: grasshopper.com

Monte Carlo Hypothesis Testing

Knowledge of test statistic distribution not needed!

x_1, \dots, x_n - data from $N(\mu, \sigma^2)$

$H_0 : \mu = 4$

$H_1 : \mu \neq 4$


Test statistic

$$T(x) = \frac{\bar{x} - \mu}{s/\sqrt{n}} \sim t(n-1)$$

- 1: **for** $i = 1$ to B **do**
- 2: Generate Y_1, \dots, Y_n i.i.d. from H_0 , i.e. $N(4, \sigma^2)$, σ^2 est.
- 3: Compute t_i from Y_1, \dots, Y_n
- 4: **end for**
- 5: Use t_1, \dots, t_B to construct a histogram
- 6: Use the histogram as the distribution of $T(x)$ under H_0

Monte Carlo Hypothesis Testing

```
x<-rnorm(10,4,1)
s<-var(x)
B<-10000
n<-length(x)
tsamp<-rep(NA,B)
for (i in 1:B){
  Y<-rnorm(n,4,s)
  tsamp[i]<-(mean(Y)-4)/(sd(Y)/sqrt(length(Y)))
}
hist(tsamp,breaks=50,col=gray(0.8),main="",xlab="t",
      ,ylab="",freq=FALSE,cex.axis=1.5,cex.lab=1.5)
```



Permutation Tests

- Introduced by Fisher 1930's
- Computationally expensive
- Applicable to certain types of hypothesis testing:
 - Equality of models, populations, etc.
- Two samples problem:
 - Population 1 distributed as F
 - Population 2 distributed as G

$$H_0 : F = G$$

$$H_1 : F \neq G$$

Permutation Tests: Mouse Data

Comparing two groups:

- Control - "z"
- Treatment - "y"

Data:


```
> t(mouse)
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
Group "y"  "z"  "z"  "y"  "y"  "z"  "y"  "y"  "y"  "y"  "z"  "z"
Value " 10" " 16" " 23" " 27" " 31" " 38" " 40" " 46" " 50" " 52" " 94" " 99"
      [,13] [,14] [,15] [,16]
Group "y"  "z"  "y"  "z"
Value "104" "141" "146" "197"
```

Do the values differ significantly between control and treatment groups?

Permutation Tests: Idea

IDEA: If $F = G$ then *group label does not matter*

→ permute labels and still have sample from F (or G)



Test statistic:

$$T = \text{mean}(\text{values} | \text{group} = z) - \text{mean}(\text{values} | \text{group} = y)$$

Permutation Test: Algorithm

- 1: $T(x)$ value of statistic from observed data
- 2: Create permutations g_1^*, \dots, g_B^* of group variable
If number of (all) permutations too large
→ sample B^* of them randomly *without replacement*
further B^* instead of B
- 3: Evaluate test statistic on each permutation
- 4: Estimate p-value:
One-sided test: $\hat{p} = \# \{T(X_{g_b^*}) \geq T(X)\} / B$
Two-sided test: $\hat{p} = \# \{|T(X_{g_b^*})| \geq |T(X)|\} / B$
- 5: p-value law \Rightarrow difference significant \Rightarrow reject H_0

Permutation Test: Example

Data	Group	A	B	A	B	B	A	$\bar{A} = 5/3$
	Value	1	2	2	1	1	2	$\bar{B} = 4/3$
Set 1	Group	A	A	B	A	B	B	$\bar{A} = 4/3$
	Value	1	2	2	1	1	2	$\bar{B} = 5/3$
Set 2	Group	B	B	A	B	A	A	$\bar{A} = 5/3$
	Value	1	2	2	1	1	2	$\bar{B} = 4/3$
Set 3	Group	B	B	B	A	A	A	$\bar{A} = 4/3$
	Value	1	2	2	1	1	2	$\bar{B} = 5/3$

Observed difference: $\bar{A} - \bar{B} = 5/3 - 4/3 = 1/3$


From permuted data: $\bar{A} - \bar{B} \in \{-1/3, 1/3, -1/3\}$

Difference between two groups significant?

Permutation Tests: Mouse Data


Do we reject H_0 ?

```
B=1000
stat=numeric(B)
n=dim(mouse)[1]
for(b in 1:B){
  Gb=sample(mouse$Group, n)
  stat[b]=mean(mouse$Value[Gb=='z'])-mean(mouse$
    Value[Gb=='y'])
}
stat0=mean(mouse$Value[mouse$Group=='z'])-mean(
  mouse$Value[mouse$Group=='y'])
print(c(stat0, mean(stat>stat0)))
## [1] 30.63492 0.12700
```



Resampling Methods

Observed data




\bar{X}

Sampling with
replacement



Sampling without
replacement

Resampled data



$\bar{X}_1^*, \bar{X}_2^*, \dots, \bar{X}_N^*$

Jackknife and Bootstrap

Theory *different* **but** coding *similar*

$D = (X_1, \dots, X_n)$ - data from $F(\cdot, w)$

$\hat{w} = T(D)$ - estimator

1: **for** $i = 1, \dots, B$ **do**

2: Generate

$D_i^* = (X_1^*, \dots, X_n^*)$ by sampling *with replacement*
 $\{Nonparametric\ Bootstrap, F \text{ unknown}\}$

$D_i^* = X_{[-i]}$ $\{Jackknife, F \text{ unknown}\} \rightarrow \text{see next slide}$

$D_i^* = (X_1^*, \dots, X_n^*)$ by generating from $F(\cdot, \hat{w})$ $\{Parametric\ Bootstrap, F \text{ known}\}$

3: **end for**

4: Use histogram by $T(D_1^*), \dots, T(D_B^*)$ as distribution of \hat{w}

Jackknife: Details

- Partition sample into r groups
- Group size $k \geq 1$
- For simplicity: $n = r \cdot k$
- $X_{[-i]}$ sample without group i , $i = 1, \dots, r$
 - size of $X_{[-i]}$: $n - k$
- In step i compute estimation T from $X_{[-i]}$
 - $B \leq n$ (or $B \leq r$ if $k > 1$)

Nonparametric Bootstrap: Example

Data: 1, 2, 2, 1, 3, 4, 1	$\bar{X} = 14/7 = 2$
Set 1: 1, 2, 2, 2, 3, 1, 1	$\bar{X}_1^* = 12/7$
Set 2: 1, 1, 1, 3, 3, 2, 2	$\bar{X}_2^* = 13/7$
Set 3: 4, 4, 4, 3, 1, 2, 1	$\bar{X}_3^* = 19/7$
Set 4: 2, 1, 1, 2, 1, 1, 4	$\bar{X}_4^* = 12/7$
Set 5: 3, 4, 4, 3, 2, 2, 4	$\bar{X}_5^* = 22/7$

Bootstrap sample: $\{12/7, 13/7, 19/7, 12/7, 22/7\}$

$$\text{Var} [\bar{X}^*] = \text{Var} [\{12/7, 13/7, 19/7, 12/7, 22/7\}]$$

$$\begin{aligned}\text{Quantiles: } P(\bar{X} \leq 3) &\approx 4/5 \\ P(\bar{X} \leq 2) &\approx 3/5\end{aligned}$$

IDEA: Instead of working with true (unknown) distribution underlying \bar{X} we take the observed sample as substitute

Bootstrap Confidence Intervals

Estimate $100(1 - \alpha)\%$ percentile confidence interval for w

1: Compute $T(D_1^*), \dots, T(D_B^*)$ using Bootstrap

2: *Bootstrap percentile method:*

Sort $T(D_1^*), \dots, T(D_B^*)$ in ascending order obtaining y_1, \dots, y_B

Bootstrap t-method:

Compute $y_i = (T(D_i^*) - T(D)) / (se(T(D_i^*)))$, $i = 1, \dots, B$

3: Define $A_1 = \lceil (B \cdot \alpha/2) \rceil$, $A_2 = \lfloor (B - B \cdot \alpha/2) \rfloor$

4: Confidence interval:

(y_{A_1}, y_{A_2}) {percentile method} **or**

$(T(D) - se(T(D^*)) \cdot y_{A_1}, T(D) + se(T(D^*)) \cdot y_{A_2})$ {t-method}

Statistic from observed data not in CI \rightarrow reject H_0

Uncertainty Estimation: Variance of Estimator

Bootstrap

$$\widehat{\text{Var}}[T(\cdot)] = \frac{1}{B-1} \sum_{i=1}^B \left(T(D_i^*) - \overline{T(D^*)} \right)^2$$

Jackknife ($n = B$)


$$\widehat{\text{Var}}[T(\cdot)] = \frac{1}{n(n-1)} \sum_{i=1}^n (T_i^* - J(T))^2,$$

where

$$T_i^* = nT(D) - (n-1)T(D_i^*) \quad J(T) = \frac{1}{n} \sum_{i=1}^n T_i^*$$

Bootstrap in R

```
library("boot")
stat1<-function(data,vn){
  data<-as.data.frame(data[vn,])
  res<-lm(Response~Predictor,data)
  res$coefficients[2]
}
x<-rnorm(100);data<-cbind(Predictor=x,Response=3+2*
  x+rnorm(length(x),sd=0.5))
res<-boot(data,stat1,R=1000)
print(boot.ci(res))
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
##Based on 1000 bootstrap replicates
#Intervals :
#Level      Normal              Basic
#95%      ( 1.933,  2.164 )    ( 1.935,  2.162 )
# Level      Percentile          BCa
#95%      ( 1.934,  2.161 )    ( 1.936,  2.166 )
```



Bootstrap Bias Correction

- 1: Observed data: $D = (X_1, \dots, X_n)$, estimator $\hat{w} = T(D)$
- 2: **for** $i = 1, \dots, B$ **do**
- 3: Generate $D_i^* = (X_1^*, \dots, X_n^*)$ by sampling with replacement
- 4: Calculate $T_i^* = T(D_i^*)$
- 5: **end for**
- 6: Bias corrected estimator is

$$T_1 := 2 \cdot T(D) - \frac{1}{B} \sum_{i=1}^B T_i^*$$

Jackknife also has a bias correction method

Comments: Resampling Methods

- Bootstrap-t method is more accurate than percentile
- Permutations: sampling *without* replacement
bootstrap: sampling *with* replacement
- Permutation: p-value *exact* if all permutations used
bootstrap: p-value always *approximate*
- Bootstrap may be used for a wider class of problems
- Methods do not require distributional assumptions

Summary

- Hypothesis testing
- Monte Carlo hypothesis testing
- Resampling methods
 - permutations
 - jackknife
 - nonparametric bootstrap
- Simulation methods
 - parametric bootstrap