

Basic concepts

Lecture 1a

Course leader: Oleg Sysoev



Course topics

Block 1

- Basic concepts in machine learning. Software for ML. Classification and regression
- Dimensionality reduction and model selection
- Kernel methods (SVM) and neural networks

Block 2

- Mixture models and ensemble methods

Course organization

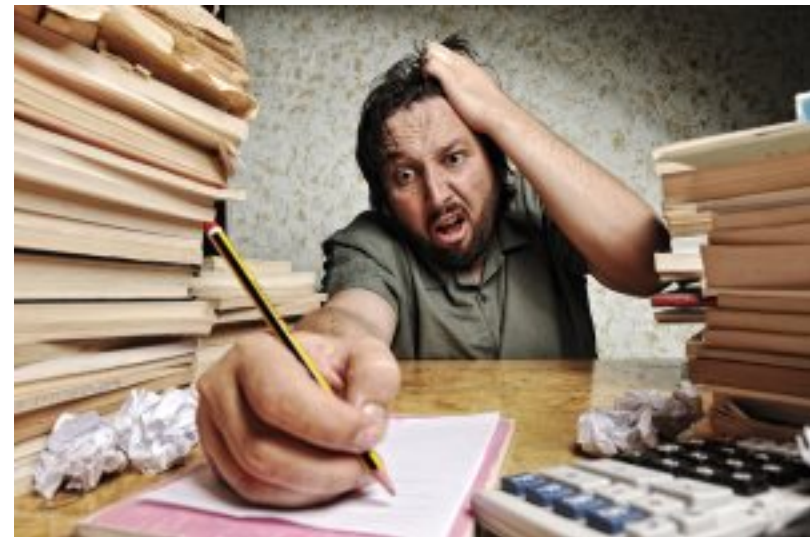
- 1 topic= 3-4 lectures (Zoom) +1 lab (2h* 3, campus)+seminar (zoom)
- Course given as
 - 732A99 (9 ECTS): Block 1+Block 2
 - 732A68 (9 ECTS): Block 1+Block 2
 - TDDE01 (6 ECTS): Block 1
- **Labs**
 - Sign-up at LISAM, **exactly 3 persons!** (otherwise group can be split)
 - Take around 8h , group report
 - Recommended strategy: one person has main responsibility for one assignment in the lab.
 - **Statement of Contribution:** describe clearly how each member contributed to the group report (which member which assignment). Without it lab is automatically failed.
 - Deadlines
 - To pass exam, **each individual needs to have experience of solving all lab tasks** → make sure to try all tasks before the exam!
 - Published a couple of days in advance – try doing before attending to the first lab session!
 - Submission via LISAM

Course organization

- Lectures
 - Available as PowerPoint or PDF, normally at LISAM
 - Write in the chat if you have questions during the lectures
- Seminars
 - Obligatory attendance of all seminars
 - Zoom
 - Speaker and opponent groups
 - Discussion of the latest lab
 - Presentation schedule will be published on LISAM (Seminar.PDF)

Course organization

- Examination
 - laboratory part + computer-based exam
- Lecture 1b is 'Basic Statistics'
- Lecture 1c is 'Introduction to R'



<http://www.swagseduction.com/wp-content/uploads/2014/11/stressful.jpg>

What is Machine Learning ?

- Machine learning is a subfield of **computer science** that evolved from the study of **pattern recognition** and computational learning theory in **artificial intelligence**.
- Machine learning explores the study and construction of **algorithms** that can **learn** from and make **predictions** on **data**. Such algorithms operate by building a model from example inputs in order to make data-driven predictions or **decisions**, rather than following strictly static program instructions.

Wikipedia (Oct 15, 2016).

Machine Learning and Statistics

- ML=**intersection** of **computer science**, **statistics** and **artificial intelligence**.
 - Related: **data mining**, **knowledge discovery** and **data science**.
- ML uses mainly **statistical (probabilistic) models** for **analyzing data**.
 - Data mining and knowledge discovery tend to use less rigorous, but often effective, algorithms.
 - ML is not a discovery of a hidden information (Data Mining)
- ML vs Statistics: ML has a **heavier focus on prediction**, and lesser on interpretation.
- ML applications often involve large sets → **computational complexity** of algorithms is important.
 - Statistics often does not care about runtime

Why probability models?

- Probability models and statistical inference provide a **framework**
- A principled **way to think** about any problem in machine learning
 - Probabilistic model \rightarrow Estimation \rightarrow Prediction
- Probabilistic models **quantify uncertainties**.
 - Deterministic answers may often be inappropriate



<http://lolnada.org/t/src/1454993210255.jpg>

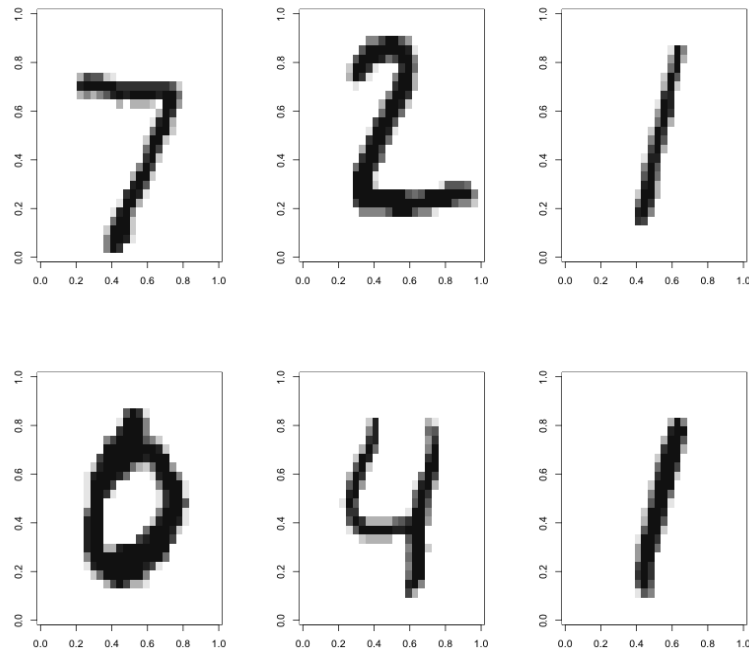
The currency exchange rate tomorrow will be 10.41!

Why probability models?

*As robotics is now moving into the open world, the issue of **uncertainty** has become a major stumbling block for the design of capable robot systems. Managing uncertainty is possibly the most important step towards robust real-world robot systems.*

from the book Probabilistic Robotics by Thrun et al.

Example: classifying handwritten digits



Example: classifying handwritten digits

Training data: 60000 images.

Test data: 10000 images.

Features: intensities (0-255, scaled to 0-1) in the $28 \times 28 = 784$ pixels as features.

Methods:

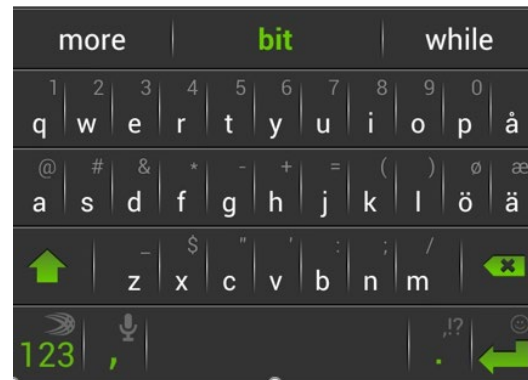
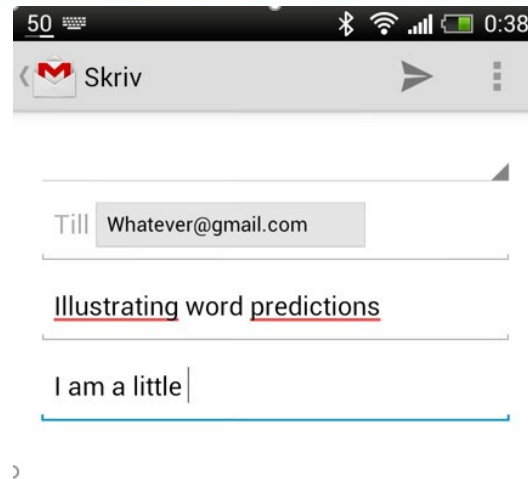
- Multinomial regression with LASSO prior
- Support vector machines
- Neural Networks (deep?)

Example: classifying handwritten digits

- Confusion matrix

		PREDICTION									
T R U T H		0	1	2	3	4	5	6	7	8	9
	0	966	0	8	1	1	7	9	2	4	6
	1	0	1121	1	1	0	2	3	13	7	7
	2	2	2	957	13	5	4	4	21	7	0
	3	0	2	9	947	0	29	1	3	12	10
	4	0	0	12	1	940	5	5	9	8	32
	5	6	1	3	19	1	816	9	1	24	9
	6	4	4	13	1	7	12	926	0	10	1
	7	1	0	9	10	2	2	0	954	5	13
	8	1	4	17	11	2	10	1	3	892	4
	9	0	1	3	6	24	5	0	22	5	927

Example: smartfone typing predictions



Example: smartfone typing predictions

- Assume a simple (Markov) model of a sentence:

$$p(w_1, \dots, w_n) = p(w_1)p(w_2|w_1) \dots p(w_n|w_{n-1})$$

- Intuition:

- $p(\text{person}|\text{crazy}) = 0.1$
- $p(\text{horse}|\text{crazy}) = 0.0001$

Highest $P(?|\text{Donald})$?


- Probability for sentence depends only on $p(w_n|w_{n-1})$
- How to compute ? Investigate a lot of data!

$$p(w_k|w_{k-1}) = \frac{\# \text{ cases } w_k \text{ follows } w_{k-1}}{\# \text{ cases } w_{k-1}}$$

- In practice, more advanced model used
 - Neural networks for ex.

Types of learning

- **Supervised learning** (classification, regression)
 - Compute parameters from data
 - Given features of a new object, predict target (generalize beyond seen training data)
 - **Classification** (Y =categorical), **Regression** (Y =continuous)
- Most of ML models: Neural Nets, Decision Trees, Support Vector Machines, Bayesian nets



Types of learning

- Unsupervised learning (→ Data Mining)
 - No target
 - Aim is to extract interesting information about
 - Relations of parameters to each other
 - Grouping of objects

Ex: clustering, density estimation, association analysis

$x1 \leftrightarrow x2 \leftrightarrow x3 \dots$

Types of learning

- **Semi-supervised**: targets are known only for some observations.
- **Active learning**. Strategies for deciding which observations to label
- **Reinforcement learning**. Find suitable actions to maximize the reward. True targets are discovered by trial and error.
- **Transfer learning**: use knowledge from some domain to train better models in a similar domain

Basic ML ingredients

- **Data** T : observations (cases)

- Features x_1, \dots, x_p
- Targets y_1, \dots, y_r

Case	x_1	x_2	y
1			
2			
...			

- **Mathematical Model** $P(x | w_1, \dots, w_k)$ or $P(y | x, w_1, \dots, w_k)$
 - Example: Linear regression $p(y | x, w) = N(w_0 + w_1 x, \sigma^2)$
- **Learning algorithm** (data \rightarrow get parameters \hat{w} or $p(w | D)$)
 - Maximum likelihood, Bayesian estimation...
- **Prediction** of new data x_* by using the fitted model

Types of data sets

- **Training data** (training set T): used for learning the model

- Supervised learning: w_i in $P(y|\mathbf{x}, w_1, \dots, w_k)$ estimated using T

X	Y
1.1	M
2.3	F

- **Test data** (test set T*): used for predictions


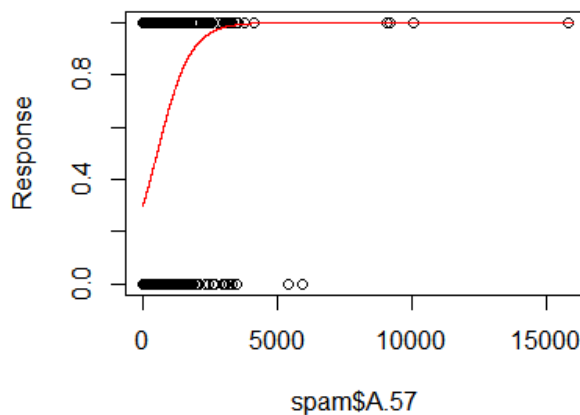
- Supervised learning: estimate $p(y_*)$ or \hat{y}_* for new \mathbf{x}_*

X	Y
1.3	?
2.9	?

Logistic regression


- **Data** $y_i \in \{Spam, Not\ Spam\}$, $x_i = \#of\ a\ word$
- **Model**: $p(y = Spam|w, x) = \frac{1}{1+e^{-w_0-w_1x}}$
- **Learning algorithm**: maximum likelihood
- **Prediction** : $p(spam) = p(Y = spam|x_*)$

We can also make point predictions
-how?



K-nearest neighbor model

- Can be classification or regression
- Basic idea:
 - For given x_* , find K nearest observations
 - Classification: majority voting
 - Regression: compute mean
- K is called **hyperparameter**



K-nearest neighbor algorithm

Data: Training data $\{\mathbf{x}_i, y_i\}_{i=1}^n$ and test input \mathbf{x}_\star

Result: Predicted test output $\hat{y}(\mathbf{x}_\star)$

- 1 Compute the distances $\|\mathbf{x}_i - \mathbf{x}_\star\|_2$ for all training data points $i = 1, \dots, n$
- 2 Let $\mathcal{N}_\star = \{i : \mathbf{x}_i \text{ is one of the } k \text{ data points closest to } \mathbf{x}_\star\}$
- 3 Compute the prediction $\hat{y}(\mathbf{x}_\star)$ as



$$\hat{y}(\mathbf{x}_\star) = \begin{cases} \text{Average}\{y_j : j \in \mathcal{N}_\star\} & \text{(Regression problems)} \\ \text{MajorityVote}\{y_j : j \in \mathcal{N}_\star\} & \text{(Classification problems)} \end{cases}$$

K-nearest neighbor model

- **Data** $T = \{(\mathbf{x}_i, y_i), i = 1, \dots, n\}$
- **Model**: W same size as T
- **Learning algorithm**: Set $W=T$, compute distances in W
- **Prediction**:
 - $y_* = \frac{1}{|N_*|} \sum_{i \in N_*} y_i$
 - $y_* = \text{MajorityVote}_{j \in N_*}(y_j)$

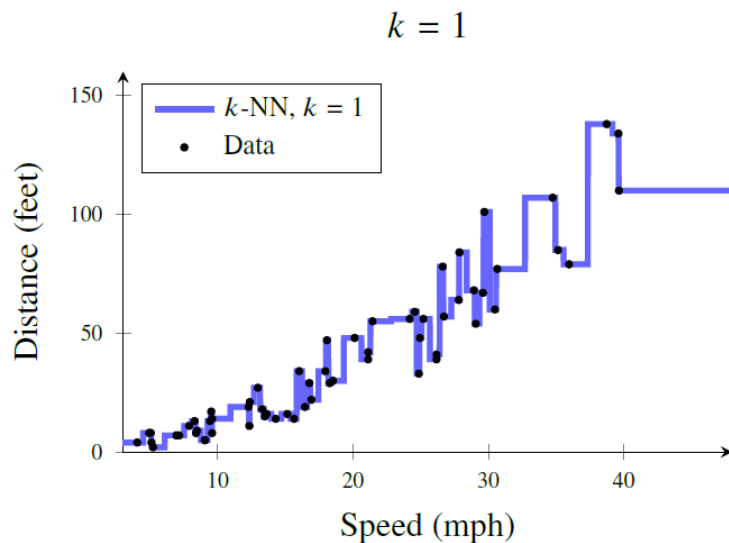
K-nearest neighbor example

- Classification
 - Music data, x_1 =song length, x_2 =a signal processing characteristic




K-nearest neighbor example

- Regression
 - Car data: x1 speed when brake signal given, x2 distance until full stop





How to choose K ?



How to define distance?




K-nearest neighbor model

- **Feature preprocessing:** scaling
 - Important, for ex when defining distance
 - Usual preprocessing: $x'_{ij} = \frac{x_{ij} - \text{mean}(x_{ij} | i=1, \dots, n)}{\text{std}(x_{ij} | i=1, \dots, n)}$






Overfitting

- Which model feels appropriate?



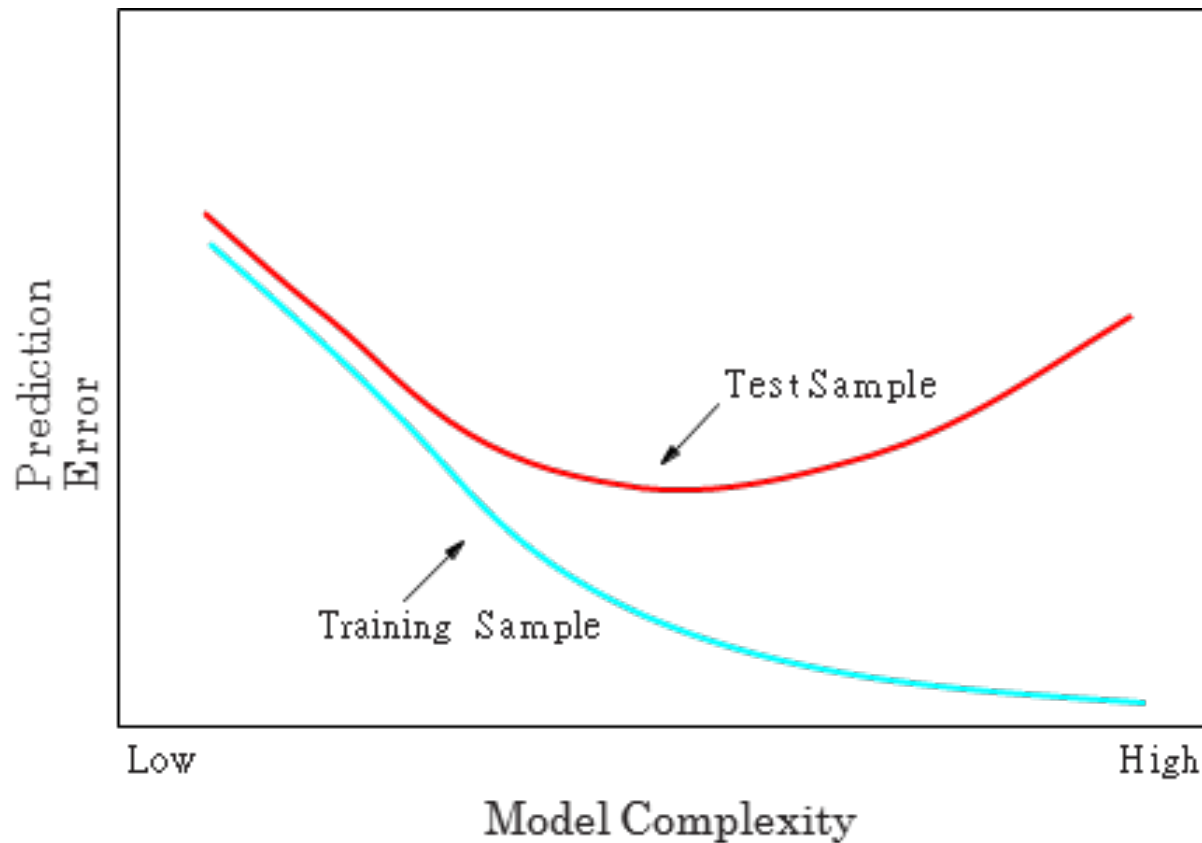
Overfitting

Now new data from the same process




Overfitting

- Observed:



Model selection

- Given several models M_1, \dots, M_m
- Divide data set into **training** and **test** data




- Fit models M_i to training data → get parameter values
- Use fitted models to predict test data and compare **test errors** $R(M_1), \dots, R(M_m)$
- Model with lowest prediction error is best

Comment:

- Approach works well for moderate/large data

Holdout method

Divide into training, validation and test sets



- Choose proportions in some way
- Test set is used to test a performance on a new data

Holdout in R

- How to partition into train/test?
 - Use `set.seed(12345)` in the labs to get identical results

```
n=dim(data)[1]
set.seed(12345)
id=sample(1:n, floor(n*0.7))
train=data[id,]
test=data[-id,]
```

- How to partition into train/valid/test?

```
n=dim(data)[1]
set.seed(12345)
id=sample(1:n, floor(n*0.4))
train=data[id,]

id1=setdiff(1:n, id)
set.seed(12345)
id2=sample(id1, floor(n*0.3))
valid=data[id2,]

id3=setdiff(id1,id2)
test=data[id3,]
```


Typical error functions

- Regression, **MSE** :

$$R(Y, \hat{Y}) = \frac{1}{N} \sum_{i=1}^N (Y_i - \hat{Y}_i)^2$$

- Classification, **misclassification rate**



$$R(Y, \hat{Y}) = \frac{1}{N} \sum_{i=1}^N I(Y_i \neq \hat{Y}_i)$$

- Classification, cross-entropy for M classes C_1, \dots, C_M :

$$R(Y, \hat{p}(Y)) = - \sum_{i=1}^N \sum_{m=1}^M I(Y_i = C_m) \log \hat{p}(Y_i = C_m)$$

Model types




- Parametric models
 - Have certain number of parameters independently of the size of training data
 - Assumption about of the data distribution
 - Ex: logistic regression
- Nonparametric models
 - Number of parameters (complexity) changes with training data
 - Example: K-NN classifier



Curse of dimensionality

- Given data T :
 - Features x_1, \dots, x_p
 - Targets y_1, \dots, y_r
- When p increases models using “proximity” measures work badly
- **Curse of dimensionality**: A point has no “near neighbors” in high dimensions → using class labels of a neighbor can be misleading
 - Distance-based methods affected

Curse of dimensionality



Curse of dimensionality

- Hopeless? No!
- Real data normally has much lower effective dimension
 - Dimensionality reduction techniques
- Smoothness assumption
 - small change in one of x 's should lead to small change in $y \rightarrow$ interpolation